

---

## ✓ Numpy 1

---

### ✓ Content

- Introduction to DAV
  - Python Lists vs Numpy Array
    - Importing Numpy
    - Why use Numpy?
  - Dimension & Shape
  - Type Conversion
  - Indexing
  - Slicing
  - NPS use case
- 

### ✓ Introduction to DAV (Data Analysis and Visualization) Module

With this lecture, we're starting the DAV module.

It will contain 3 sections -

1. DAV-1: Python Libraries
    - Numpy
    - Pandas
    - Matplotlib & Seaborn
  2. DAV-2: Probability Statistics
  3. DAV-3: Hypothesis Testing
- 

### ✓ Python Lists vs Numpy Arrays

#### ✓ Homogeneity of data

So far, we've been working with Python lists, that can have **heterogenous data**.

```
a = [1, 2, 3, "Michael", True]
a
[1, 2, 3, 'Michael', True]
```

Because of this heterogeneity, in Python lists, the data elements are not stored together in the memory (RAM).

- Each element is stored in a different location.
- Only the address of each of the element will be stored together.
- So, a list is actually just referencing to these different locations, in order to access the actual element.


On the other hand, Numpy only stores **homogenous data**, i.e. a numpy array cannot contain mixed data types.

It will either

- ONLY contain integers
- ONLY contain floats
- ONLY contain characters

... and so on.

Because of this, we can now store these different data items together, as they are of the same type.



## Speed

Programming languages can also be slow or fast.

In fact,

- Java is a decently fast language.
- Python is a slow language.
- C, one of the earliest available languages, is super fast.

This is because C has concepts like memory allocation, pointers, etc.

### How is this possible?

With Numpy, though we will be writing our code using Python, but behind the scene, all the code is written in the **C programming language**, to make it faster.

Because of this, a Numpy Array will be significantly faster than a Python List in performing the same operation.

This is very important to us, because in data science, we deal with huge amount of data.

## ✓ Properties

### • In-built Functions

- For a Python list `a`, we had in-built functions like `.sum(a)`, etc.
- For NumPy arrays also, we will have such in-built functions.

### • Slicing

- Recall that we were able to perform list slicing.
- All of that is still applicable here.

## ✓ Importing Numpy

Recall how we used to import a module/library in Python.

- In order to use Python Lists, we do not need to import anything extra.
- However to use Numpy Arrays, we need to import it into our environment, as it is a Library.

Generally, we do so while using the alias `np`.

```
import numpy as np
```

### Note:

- In this terminal, we will already have numpy installed as we are working on Google Colab
- However, when working on an environment that does not have it installed, you'll have to install it the first time working.
- This can be done with the command: `!pip install numpy`

## ✓ Why use Numpy? - Time Comparison

Suppose you are given a list of numbers. You have to find the square of each number and store it in the original list.

```
a = [1,2,3,4,5]
```

```
type(a)  
list
```

The basic approach here would be to iterate over the list and square each element.

```
res = [i**2 for i in a]  
print(res)  
[1, 4, 9, 16, 25]
```

Let's try the same operation with Numpy.

To do so, first of all we need to define the Numpy array.

We can convert any list `a` into a Numpy array using the `array()` function.

```
b = np.array(a)  
b  
array([1, 2, 3, 4, 5])
```

```
type(b)  
numpy.ndarray
```

- `nd` in `numpy.ndarray` stands for **n-dimensional**

Now, how can we get the square of each element in the same Numpy array?

```
b**2  
array([ 1,  4,  9, 16, 25])
```

**The biggest benefit of Numpy is that it supports element-wise operation.**

Notice how easy and clean is the syntax.

But is the clean syntax and ease in writing the only benefit we are getting here?

- To understand this, let's measure the time for these operations.
- We will use `%timeit`.

```
l = range(1000000)
```

```
%timeit [i**2 for i in l]  
343 ms ± 21.6 ms per loop (mean ± std. dev. of 7 runs, 1 loop each)
```

**It took approx 300 ms per loop to iterate and square all elements from 0 to 999,999**

Let's perform the same operation using Numpy arrays -

- We will use `np.array()` method for this.
- We can perform element wise operation using numpy.

```
l = np.array(range(1000000))
```

```
%timeit l**2  
778 µs ± 100 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)
```

Notice that it only took 900  $\mu$ s per loop time for the numpy operation.

**What is the major reason behind numpy's faster computation?**

- Numpy array is densely packed in memory due to its **homogenous** type.

- Numpy functions are implemented in **C programming language**.
- Numpy is able to divide a task into multiple subtasks and process them **parallelly**.

## ✓ Dimensions and Shape

We can get the dimension of an array using the `ndim` property.

```
arr1 = np.array(range(1000000))
arr1.ndim

1
```

Numpy arrays have another property called `shape` that tells us number of elements across every dimension.

```
arr1.shape

(1000000,)
```

This means that the array `arr1` has 1000000 elements in a single dimension.

Let's take another example to understand `shape` and `ndim` better.

```
arr2 = np.array([[1, 2, 3], [4, 5, 6], [10, 11, 12]])
print(arr2)

[[ 1  2  3]
 [ 4  5  6]
 [10 11 12]]
```

What do you think will be the shape & dimension of this array?

```
arr2.ndim

2
```


```
arr2.shape

(3, 3)
```

`ndim` specifies the number of dimensions of the array i.e. 1D (1), 2D (2), 3D (3) and so on.


`shape` returns the exact shape in all dimensions, that is (3,3) which implies 3 in axis 0 and 3 in axis 1.

### 1D array




shape: (4,)

### 2D array



shape: (2, 3)

### 3D array



shape: (4, 3, 2)

## ✓ np.arange()

Let's create some sequences in Numpy.

We can pass **starting** point, **ending** point (not included in the array) and **step-size**.

**Syntax:**

- `arange(start, end, step)`

```
arr2 = np.arange(1, 5)
arr2
array([1, 2, 3, 4])
```

```
arr2_step = np.arange(1, 5, 2)
arr2_step
array([1, 3])
```

`np.arange()` behaves in the same way as `range()` function.

**But then why not call it `np.range`?**

- In `np.arange()`, we can pass a **floating point number** as **step-size**.

```
arr3 = np.arange(1, 5, 0.5)
arr3
array([1. , 1.5, 2. , 2.5, 3. , 3.5, 4. , 4.5])
```

---

## ✓ Type Conversion in Numpy Arrays

- For this, let's pass a **float** as one of the values in a **numpy array**.

```
arr4 = np.array([1, 2, 3, 4])
arr4
array([1, 2, 3, 4])
```

```
arr4 = np.array([1, 2, 3, 4.0])
arr4
array([1., 2., 3., 4.])
```

- Notice that **int is raised to float**
- Because a numpy array can only store **homogenous data** i.e. values of one data type.

Similarly, what will happen when we run the following code? Will it give an error?

```
np.array(["Harry Potter", 1, 2, 3])
array(['Harry Potter', '1', '2', '3'], dtype='<U21')
```

No. It will convert all elements of the array to `char` type.

There's a `dtype` parameter in the `np.array()` function.

**What if we set the `dtype` of array containing integer values to `float`?**

```
arr5 = np.array([1, 2, 3, 4])
arr5
array([1, 2, 3, 4])
```

```
arr5 = np.array([1, 2, 3, 4], dtype="float")
arr5

array([1., 2., 3., 4.])
```

**Question:** What will happen in the following code?

```
np.array(["Shivank", "Bipin", "Ritwik"], dtype=float)

-----
ValueError                                Traceback (most recent call last)
<ipython-input-26-bdb627c3c07e> in <cell line: 1>()
----> 1 np.array(["Shivank", "Bipin", "Ritwik"], dtype=float)

ValueError: could not convert string to float: 'Shivank'
```

SEARCH STACK OVERFLOW

Since it is not possible to convert strings of alphabets to floats, it will naturally return an Error.

We can also convert the data type with the `astype()` method.

```
arr = np.array([10, 20, 30, 40, 50])
arr

array([10, 20, 30, 40, 50])

arr = arr.astype('float64')
print(arr)

[10. 20. 30. 40. 50.]
```

---

## ✓ Indexing

- Similar to Python lists

```
m1 = np.arange(12)
m1

array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
```

```
m1[0] # gives first element of array

0
```

```
m1[-1] # negative indexing in numpy array

11
```

You can also use list of indexes in numpy.

```
m1 = np.array([100,200,300,400,500,600])

m1[[2,3,4,1,2,2]]

array([300, 400, 500, 200, 300, 300])
```

Did you notice how single index can be repeated multiple times when giving list of indexes?

**Note:**

- If you want to extract multiple indices, you need to use two sets of square brackets `[[ ]]`
  - Otherwise, you will get an error.
- Because it is only expecting a single index.
- For multiple indices, you need to pass them as a list.

```
m1[2,3,4,1,2,2]
```

```
-----
IndexError                                Traceback (most recent call last)
<ipython-input-34-0ec34089038e> in <cell line: 1>()
----> 1 m1[2,3,4,1,2,2]

IndexError: too many indices for array: array is 1-dimensional, but 6 were indexed
```

SEARCH STACK OVERFLOW

## ✓ Slicing

- Similar to Python lists

```
m1 = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10])
m1

array([ 1,  2,  3,  4,  5,  6,  7,  8,  9, 10])
```

```
m1[:5]

array([1, 2, 3, 4, 5])
```

**Question:** What'll be output of `arr[-5:-1]` ?

```
m1[-5:-1]

array([6, 7, 8, 9])
```

**Question:** What'll be the output for `arr[-5:-1: -1]` ?

```
m1[-5: -1: -1]

array([], dtype=int64)
```

## ✓ Fancy Indexing (Masking)

- Numpy arrays can be indexed with boolean arrays (masks).
- This method is called **fancy indexing** or **masking**.

What would happen if we do this?

```
m1 = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10])
m1 < 6

array([ True,  True,  True,  True,  True, False, False, False, False,
        False])
```

**Comparison operation also happens on each element.**

- All the values before 6 return `True`
- All the values after 6 return `False`

**Question:** What will be the output of the following?

```
m1[[True, True, True, True, True, False, False, False, False, False]]

array([1, 2, 3, 4, 5])
```

Notice that we are passing a list of indices.

- For every instance of `True`, it will print the corresponding index.

- Conversely, for every `False`, it will skip the corresponding index, and not print it.

So, this becomes a **filter** of sorts.

Now, let's use this to filter or mask values from our array.

**Condition will be passed instead of indices and slice ranges.**

```
m1[m1 < 6]
array([1, 2, 3, 4, 5])
```

This is known as **Fancy Indexing** in Numpy.

**Question:** How can we filter/mask even values from our array?

```
m1[m1%2 == 0]
array([ 2,  4,  6,  8, 10])
```

## ✓ Use Case: NPS (Net Promoter Score)


Imagine you are a Data Analyst @ Airbnb

You've been asked to analyze user survey data and report NPS to the management.

But, what exactly is NPS?

Have you all seen that every month, you get a survey form from Scaler?

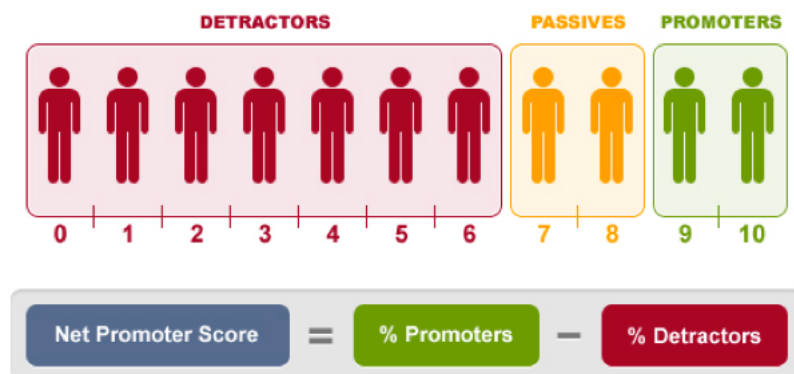
- This form asks you to fill in feedback regarding how you are liking the services of Scaler in terms of a numerical score.
- This is known as the **Likelihood to Recommend Survey**.
- It is widely used by different companies and service providers to evaluate their performance and customer satisfaction.



- Responses are given a scale ranging from 0–10,
  - with 0 labeled with "Not at all likely," and
  - 10 labeled with "Extremely likely."

Based on this, we calculate the **Net Promoter Score**.

## ✓ How to calculate NPS score?



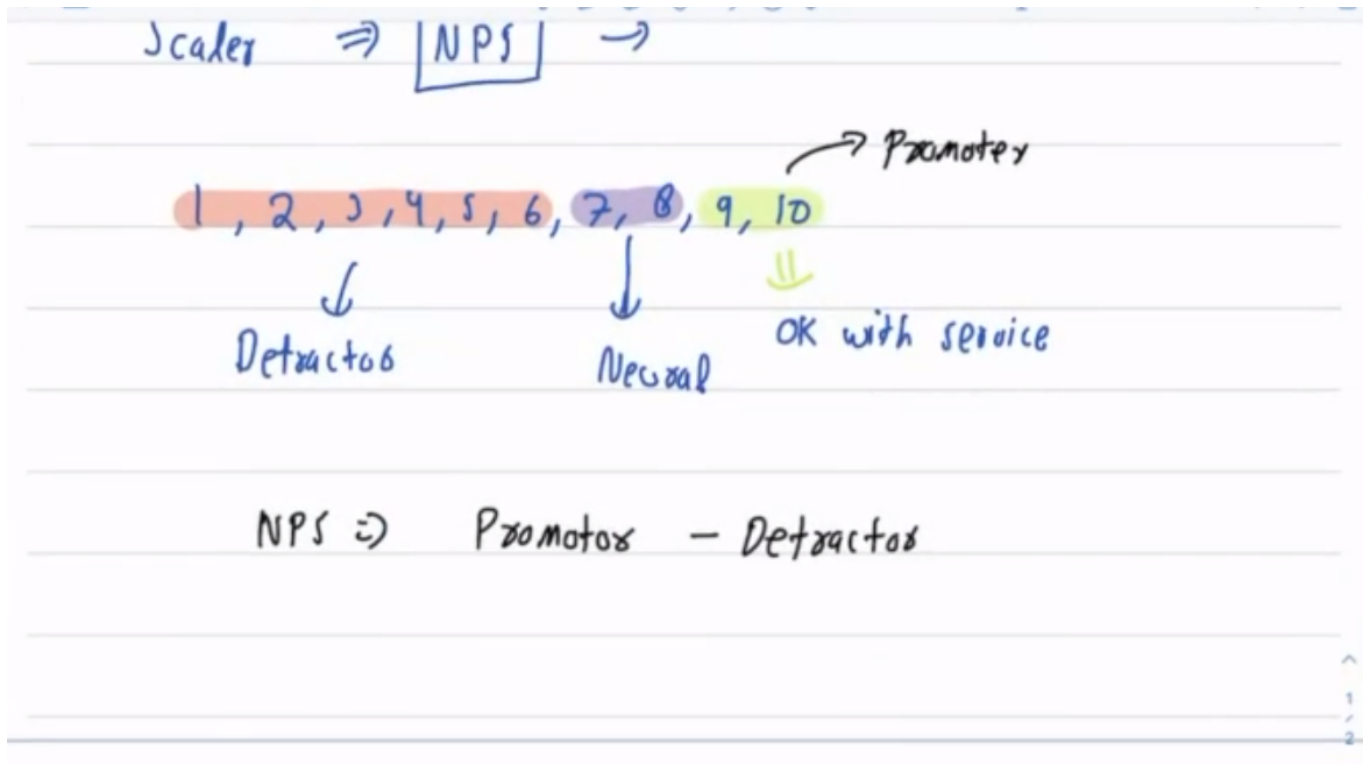
We label our responses into 3 categories:

- **Detractors:** Respondents with a score of 0-6
- **Passive:** Respondents with a score of 7-8



- **Promoters:** Respondents with a score of 9-10.

Net Promoter score = % Promoters - % Detractors.



## Range of NPS

- If all people are promoters (rated 9-10), we get 100 NPS
- Conversely, if all people are detractors (rated 0-6), we get -100 NPS
- Also, if all people are neutral (rated 7-8), we get a 0 NPS

Therefore, the range of NPS lies between  $[-100, 100]$

Generally, each company targets to get at least a threshold NPS.

- For Scaler, this is a score of 70.
- This means that if  $NPS > 70$ , it is great performance of the company.

Naturally, this varies from business to business.

## How is NPS helpful?

Why would we want to analyse the survey data for NPS?

NPS helps a brand in gauging its brand value and sentiment in the market.

- Promoters are highly likely to recommend your product or service. Hence, bringing in more business.
- whereas, Detractors are likely to recommend against your product or service's usage. Hence, bringing the business down.

These insights can help business make customer oriented decision along with product improvisation.

**Two third of Fortune 500 companies use NPS.**

Even at Scaler, every month, we randomly reach out to our learners over a call, and try to understand,

- How is the overall experience for them?
- What are some things that they like?
- What do they don't like?

Based on the feedback received, sometimes we end up getting really good insights, and tackle them.

This will help improve the next month's NPS.

## ✓ NPS Problem

Let's first look at the data we have gathered.

**Dataset:** <https://drive.google.com/file/d/1c0ClC8SrPwJq5rrkyMKyPn80nyHcFikK/view?usp=sharing>



```
7
10
5
9
9
4
7
9
9
9
10
4
9
9
7
10
5
9
7
9
9
9
1
8
5
9
9
10
10
```

Downloading the dataset -

```
!pip install --upgrade gdown
!gdown 1c0ClC8SrPwJq5rrkyMKyPn80nyHcFikK

Requirement already satisfied: gdown in /usr/local/lib/python3.10/dist-packages (4.7.3)
Collecting gdown
  Downloading gdown-5.1.0-py3-none-any.whl (17 kB)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from gdown) (4.12.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from gdown) (3.13.1)
Requirement already satisfied: requests[socks] in /usr/local/lib/python3.10/dist-packages (from gdown) (2.31.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from gdown) (4.66.1)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4->gdown) (2.
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests[socks]
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown) (3.
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests[socks]->gdown)
Requirement already satisfied: PySocks!=1.5.7,>=1.5.6 in /usr/local/lib/python3.10/dist-packages (from requests[socks]->
Installing collected packages: gdown
  Attempting uninstall: gdown
    Found existing installation: gdown 4.7.3
    Uninstalling gdown-4.7.3:
      Successfully uninstalled gdown-4.7.3
  Successfully installed gdown-5.1.0
Downloading...
From: https://drive.google.com/uc?id=1c0ClC8SrPwJq5rrkyMKyPn80nyHcFikK
To: /content/survey.txt
100% 2.55k/2.55k [00:00<00:00, 9.56MB/s]
```

Loading the data -

- For this we will use the `.loadtxt()` function
- We provide file name along with the dtype of data that we want to load.
- Documentation: <https://numpy.org/doc/stable/reference/generated/numpy.loadtxt.html>

```
score = np.loadtxt('survey.txt', dtype='int')
```

Let's check the type of this data variable `score` -

```
type(score)

numpy.ndarray
```

Let's see what the data looks like -

```
score[:5]

array([ 7, 10,  5,  9,  9])
```

Let's check the number of responses -

```
score.shape  
(1167,)
```

There are a total of 1167 responses for the LTR survey.

Now, let's calculate NPS using these response.

**NPS = % Promoters - % Detractors**

In order to calculate NPS, we need to calculate two things:

- % Promoters
- % Detractors

In order to calculate % Promoters and % Detractors, we need to get the count of promoter as well as detractor.

**Question:** How can we get the count of Promoter/Detractor ?

- We can do so by using fancy indexing (masking).

Let's get the count of promoter and detractors -

Detractors have a score  $\leq 6$

```
detractors = score[score <= 6]
```

# Number of detractors -

```
num_detractors = len(detractors)  
num_detractors  
332
```

Promoters have a score  $\geq 9$

```
promoters = score[score >= 9]
```

# Number of promoters -

```
num_promoters = len(promoters)  
num_promoters  
609
```

```
total = len(score)  
total  
1167
```

# % of detractors -

```
percentage_detractors = (num_detractors/total) * 100  
percentage_detractors  
28.449014567266495
```

# % of promoters -

```
percentage_promoters = (num_promoters/total) * 100  
percentage_promoters  
52.185089974293064
```

```
nps = percentage_promoters - percentage_detractors  
nps  
23.73607540702657
```

# Rounding off upto 2 decimal places -

```
np.round(nps, 2)
```

➡ 23.74

---